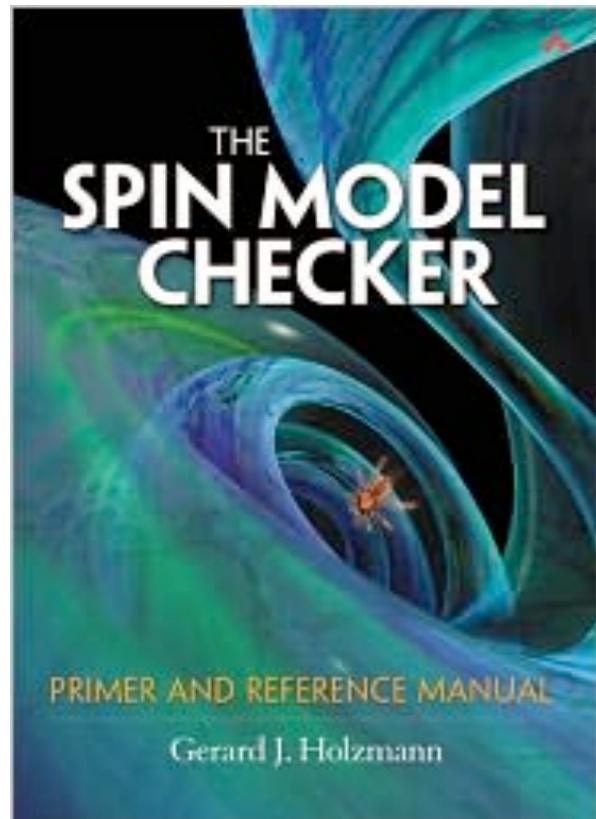


# Sumofhellos mis-using mutexes

- Let's produce a version of sumofhellos that uses mutexes in an incorrect manner. (**bad\_sumofhellos.pml**).
  - This is the advantage of a modelling language over a program language: We can say “make it so!” (to some extent, at least).
- As before, we can run simulations of this, and we observe failure.
- Again, we can build the state diagram and check every possible path through it, so determining the outcome of every possible interleaving
  - `$> spin -run bad_sumofhellos.pml`
- In this case we not only see an error indication, but a “trail” file has been created, which we can run with
  - `$> spin -p -k bad_sumofhellos.pml.trail bad_sumofhellos.pml`



# Book Reference



**The SPIN Model Checker: Primer and Reference Manual**  
By Gerald Holtzmann  
Addison-Wesley, May 2011 ISBN 0321773713  
\$65 [Paperback]

[The spin model checker : primer and reference manual /](#)  
Gerald J. Holzmann

Holzmann, Gerard J.

Printed Book | 2004

Available at Hamilton Lower (S-LEN 620 P496) plus 3 more [see all](#)

Click and Collect

Additional actions:



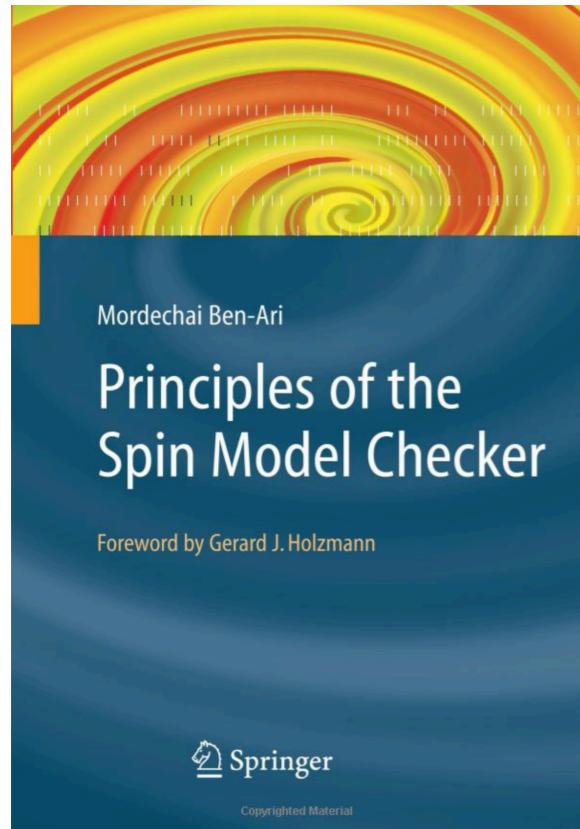
Items  only show available [close](#)

| Location (click for map) | Shelfmark        | Status |
|--------------------------|------------------|--------|
| Hamilton Lower           | S-LEN 620 P496   | IN     |
| Hamilton Lower           | S-LEN 620 P496;2 | IN     |
| Hamilton Lower           | S-LEN 620 P496;3 | IN     |
| Hamilton Upper           | 620 P496;1       | IN     |



Trinity College Dublin  
Coláiste na Tríonóide, Baile Átha Cliath  
The University of Dublin

# Book Reference



**Principles of the Spin Model Checker 2nd Edition**  
By Mordechai Ben-Ari  
Springer, 2008 ISBN 978-1-84628-769-5  
\$63 [Paperback]

**Principles of the Spin model checker / Mordechai Ben-Ari**  
Ben-Ari, M., 1948-

Printed Book | 2008

Not Available at Hamilton Lower (DUE 16-04-21) [see all](#)

Click and Collect

Additional actions:



## Items

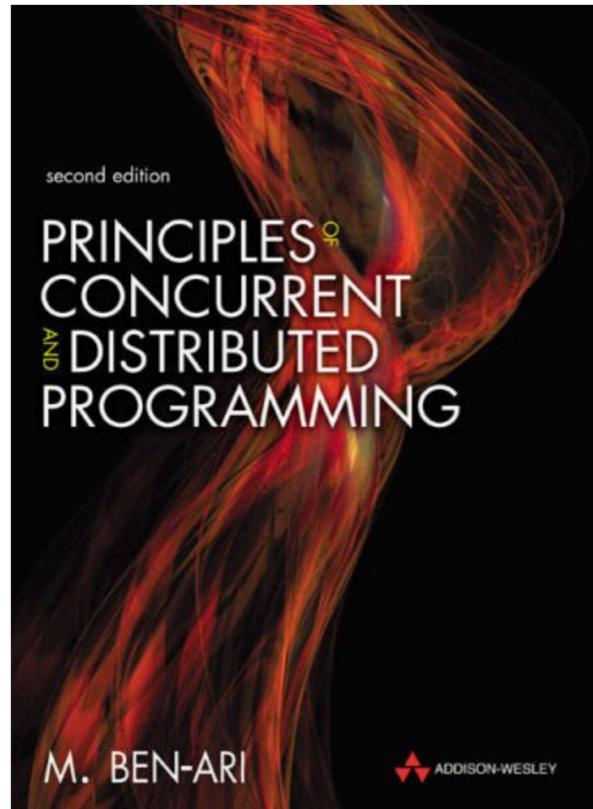
[close](#)

| Location (click for map) | Shelfmark           | Status       |
|--------------------------|---------------------|--------------|
| Hamilton Lower           | S-LEN 500.164 P83;1 | DUE 16-04-21 |
| Hamilton Upper           | 500.164 P83         | DUE 16-04-21 |



Trinity College Dublin  
Coláiste na Tríonóide, Baile Átha Cliath  
The University of Dublin

# Book Reference



**Principles of Concurrent and Distributed Programming**  
2nd Edition  
By Mordechai Ben-Ari  
Addison-Wesley, 2006 ISBN 978-0-321-31283-9  
\$120 [Paperback]

[Principles of concurrent and distributed programming /](#)

M. Ben-Ari

Ben-Ari, M., 1948-

Printed Book | 2006

Available at Hamilton Lower (S-LEN 500.164 N093\*1) plus 4+ more [see all](#)

Click and Collect

Additional actions:



[Principles of concurrent and distributed programming /](#)

M. Ben-Ari

Ben-Ari, M., 1948-

Printed Book | 1990

Available at Santry Stacks (PL-157-209) [see all](#)

Items

only show available

| Location (click for map) | Shelfmark              | Status |
|--------------------------|------------------------|--------|
| Hamilton Lower           | S-LEN 500.164 N093*1   | IN     |
| Hamilton Lower           | S-LEN 500.164 N093*1;1 | IN     |
| Hamilton Lower           | S-LEN 500.164 N093*1;2 | IN     |
| Hamilton Lower           | S-LEN 500.164 N093*1;3 | IN     |



Trinity College Dublin  
Coláiste na Tríonóide, Baile Átha Cliath  
The University of Dublin

# Mutual Exclusion (without pthreads!)

- Early challenge in concurrency was to provide mutual exclusion
  - assuming that the only atomic operations were read and write
- First algorithm, for two processes only, described in 1965 by Edsger Dijkstra
  - He attributed it to T.J Dekker in 1962.
- It does not assume that two processes alternate access
  - It allow each process to return to the critical region as many times as it wants before the other one gets access.



# Dekker's Algorithm (I)

```
bool turn;
bool flag[2];
byte count;

active [2] proctype mutex()
{
    pid i, j;

    i = _pid;
    j = 1 - _pid;

again:

    < ACQUIRE "LOCK">

    count++;
    assert(count == 1); /* critical section */
    count--;

    <RELEASE "LOCK">

    goto again
}
```

## ● Overview:

- variable **count** is the global shared variable
- variable **i** refers to thread's own id, while **j** is that of the other thread.
- Body is an infinite loop
  - defined using label **again**:
  - and a **goto** statement !



# Dekker's Algorithm (II)

```
flag[i] = true;  
do  
:: flag[j] ->  
    if  
        :: turn == j ->  
            flag[i] = false;  
            !(turn == j);  
            flag[i] = true  
        :: else  
            fi  
    :: else ->  
        break  
od;
```

- Acquiring the lock:

- set own **flag** to say "going in!"
- while other **flag** is true:
  - if other's **turn**:  
turn own **flag** off, wait for turn, turn **flag** on again
- When we exit, the other's **flag** is off and it is our **turn**.



# Dekker's Algorithm (III)

```
turn = j;  
flag[i] = false;
```

- Releasing the lock:
  - It is their **turn**
  - Our **flag** is now false.



# Checking Dekker

- We can run it: `spin -u1000 dekker.pml`
- We can verify it `spin -a dekker.pml ; cc -o pan pan.c ; ./pan`
- ???
  - SPIN takes a Promela model and compiles it into a C program that builds and analyses that model.
  - The executable is usually called “pan”
    - short for “Protocol Analyser” - the first application area for Promela/SPIN
    - splitting `spin -run model.pml` into `spin -a model.pml` and then using `pan` allows a lot of flexibility in the kinds of analyses that can be done.



# Improvement (Doran&Thomas 1980)

```
flag[i] = true;  
if  
:: flag[j] ->  
    if  
        :: turn == j ->  
            flag[i] = false;  
            !(turn == j);  
            flag[i] = true  
        :: else  
            fi;  
            (!flag[j]);  
    :: else ->  
        fi;
```

- Is top-level loop necessary?
  - Hard to say without careful analysis
- Acquiring the lock:
  - set own **flag** to say "going in!"
  - if other **flag** is true:
    - if other's **turn**:  
turn own **flag** off, wait for **turn**, turn **flag** on again
    - then, wait for their **flag** to turn off.
      - Not in SPIN/Promela book !
  - When we exit, the other's **flag** is off and it is our **turn**.



```

byte count;
byte x, y, z;
active [2] proctype user()
{ byte me = _pid + 1; /* me is 1 or 2 */
L1: x = me;
L2: if
  :: (y != 0 && y != me) -> goto L1
  :: (y == 0 || y == me)
  fi;
L3: z = me;
L4: if
  :: (x != me) -> goto L1
  :: (x == me)
  fi;
L5: y = me;
L6: if
  :: (z != me) -> goto L1
  :: (z == me)
  fi;

L7: progress:/* success: enter critical section */
  count++;
  assert(count == 1);
  count--;
  goto L1
}

```

# Something Broken

- A major computer manufacturer recommended this....
- SPIN very quickly finds a counter-example!

`spin -run mutex_flaw.pml`



# Shrinking SPIN's counter-example

- Failure is noted and a .trail file created
  - the "depth" is how many steps were taken to failure
  - We can "run" the trail: `spin -p -t mutex_flaw.pml`
- Two ways to find a shorter counter-example
  - Keep track of depth and iteratively search for minimal depth  
`cc -DREACH -o pan pan.c ; ./pan -i -mDD` where DD is known depth.
  - Use breadth-first search instead (default is depth-first)  
`cc -DBFS -o pan pan.c ; ./pan`



# The Closed World Assumption

- A Model Checker seeks to establish the truth of something by trying exhaustively and failing to demonstrate the ***negation*** of it.
  - e.g. to show that mutex holds, it tries to demonstrate (by counterexample) that it doesn't hold. If it fails to do so, it is taken as proof that the mutex holds.
- This only corresponds to what we normally understand to be “proof” if it is assumed that the model checker knows everything (i.e. can prove everything provable) about the system -- it's called the “Closed World Assumption.”
- A property deemed true, or discovered to be false, by a model-checker, may not be so in the real world.
  - This typically arises because a model only covers some aspects of a systems behaviour, and other aspects, not modelled, may also have an impact on behaviour.



# Promela Verification Constructs

- So far, we have seen that SPIN does some standard verification checks by default
  - e.g. Deadlock Freedom
- There are a number of other common cases that can be handled without using Linear Temporal Logic (LTL)
  - Basic Assertions
  - Special Labels:
    - End-States
    - Progress-States



# Promela – Basic Assertion

- Format: assert(<expression>)
  - where <expression> must evaluate to true or
    - simulation will be aborted or
    - verification will fail
- Note: an assert is evaluated during both simulation and verification.
- Already seen its use in **sumofhellos.pml**, for example



# Promela Labels

- Promela allows labels to be attached to statements
  - (e.g. see `dekker.pml` or `doran.pml`)
- By using labels with a specific **prefix**, we can trigger certain kinds of correctness checks.



# Promela – *End-State Label*

- Promela always verifies that no deadlock occurs.
- It assumes that the only valid end-states for a system are where:
  - each process is the end of its code.
- If it can show an end state where a process is not at the end of its code, it considers that an error (Deadlock).
- Sometimes, it is legitimate for a process not to end up at the end of its code. You can label such states **end...** to indicate to Promela that they are valid end points.

**end, end\_one, end00** – anything starting with **end**.



# end...: label example

- An example of a valid-end state, that is not the end of a program, is where a program that is not meant to terminate (i.e. a web-server) is waiting for client requests to arrive.
- Consider an example of two servers talking to a client that accesses them both and terminates [SMC-programs/ch04/end.pml]
  - When the client terminates, both servers are left waiting on a request
  - SPIN reports an error - but this is erroneous

```
1 /* Copyright 2007 by Moti Ben-Ari
2 under the GNU GPL; see readme.txt */
3
4 byte request = 0;
5
6 active proctype Server1() {
7   do
8     :: request == 1 ->
9       printf("Service 1\n");
10      request = 0;
11  od
12 }
13
14 active proctype Server2() {
15   do
16     :: request == 2 ->
17       printf("Service 2\n");
18      request = 0;
19  od
20 }
21
22 active proctype Client() {
23   request = 1;
24   request == 0;
25   request = 2;
26   request == 0;
27 }
```



# end...: label example

- We fix this by labelling the (start of) the `do ... od` statement in each server.
- SPIN is now happy

```
1  /* Copyright 2007 by Moti Ben-Ari
2   under the GNU GPL; see readme.txt */
3
4  byte request = 0;
5
6  active proctype Server1() {
7    end1: do
8      :: request == 1 ->
9        printf("Service 1\n");
10       request = 0;
11     od
12   }
13
14  active proctype Server2() {
15    end2: do
16      :: request == 2 ->
17        printf("Service 2\n");
18       request = 0;
19     od
20   }
21
22  active proctype Client() {
23    request = 1;
24    request == 0;
25    request = 2;
26    request == 0;
27  }
```



# Promela – Progress-State Label

- A system can have loops – cycles of sequences of states passed through infinitely often.
  - The question is, are such loops desirable or not;
    - if progress is made each time, then they are desirable loops -- progress cycles;
    - otherwise, they are undesirable *non-progress cycles*, where the system is doing something but not progressing.
  - A given loop might never get executed
    - This can happen if there is always something else that can run elsewhere, and the "scheduler" is being "unkind" to this loop
    - If the loop contains a state that should be happening repeatedly, then this is a case of **starvation**.
- You can label such states **progress**... to indicate to Promela that they should occur in any (infinite) loop.

**progress,progress\_one,progress00** – anything starting with **progress**.



# progress...: label example

- Here we another implementation of mutual exclusion
  - It is deadlock-free and only one thread is in critical region at any time.
  - A standard SPIN verification run shows no problems
- However, can one thread run at the expense of the other?

```
1 /* Copyright 2007 by Moti Ben-Ari
2 under the GNU GPL; see readme.txt */
3
4 bool wantP = false, wantQ = false;
5
6 active proctype P() {
7     do
8         :: wantP = true;
9         do
10            :: wantQ -> wantP = false; wantP = true
11            :: else -> break
12            od;
13            // critical section
14            wantP = false
15        od
16    }
17
18 active proctype Q() {
19     do
20         :: wantQ = true;
21         do
22            :: wantP -> wantQ = false; wantQ = true
23            :: else -> break
24            od;
25            // critical section
26            wantQ = false
27        od
28    }
```



# progress...: label example

- We can use a progress label
    - We label the critical section in **P()**
  - We need a special verifier:
    - Optimised for "non-progress" (NP) cycles
- cc -DNP -o pan pan.c**
- ./pan -1**

```
1  /* Copyright 2007 by Moti Ben-Ari
2  under the GNU GPL; see readme.txt */
3
4  bool wantP = false, wantQ = false;
5
6  active proctype P() {
7      do
8          :: wantP = true;
9          do
10             :: wantQ -> wantP = false; wantP = true
11             :: else -> break
12             od;
13             progressP: // critical section
14             wantP = false
15         od
16     }
17
18  active proctype Q() {
19      do
20          :: wantQ = true;
21          do
22             :: wantP -> wantQ = false; wantQ = true
23             :: else -> break
24             od;
25             // critical section
26             wantQ = false
27         od
28     }
```



# Showing infinite cycles

- SPIN has a way of showing infinite cycles

```
[:- spin -t -p SMC-5-1.pml
starting claim 2
spin: couldn't find claim 2 (ignored)
 2: proc 1 (Q:1) SMC-5-1.pml:21 (state 1) [wantQ = 1]
<<<<START OF CYCLE>>>>
 4: proc 1 (Q:1) SMC-5-1.pml:24 (state 5) [else]
 6: proc 1 (Q:1) SMC-5-1.pml:27 (state 10) [wantQ = 0]
 8: proc 1 (Q:1) SMC-5-1.pml:21 (state 1) [wantQ = 1]
spin: trail ends after 8 steps
#processes: 2
          wantP = 0
          wantQ = 1
 8: proc 1 (Q:1) SMC-5-1.pml:22 (state 7)
 8: proc 0 (P:1) SMC-5-1.pml:7 (state 11)
2 processes created
```

- it marks the start of the cycle with <<<<START OF CYCLE>>>>
- the last state shown is the same as the one just before that marker



# "never claims" ???

- When we run the liveness analysis we see that the SPIN output is suddenly all about "claims" and "never claims".

Full statespace search for:

never claim

+ (:np\_ :)

assertion violations

+ (if within scope of claim)

non-progress cycles

+ (fairness disabled)

invalid end states

- (disabled by never claim)

- What are these?



# Never Claims

- Never Claims are the mechanism used by SPIN to check LTL properties
- The state models produced by SPIN are not just simple graphs
  - They are automata, known as Büchi automata
- Büchi Automata:
  - Finite state machines
  - With criteria defined for accepting both finite and infinite languages
  - This means that accepting states can be state-loops
- LTL properties can be converted to Büchi automata
  - These can be expressed in Promela
  - They run in parallel with the Promela model defined in a `.pml` file.



# Finite State Machines/Automata (FSM/FSA)

- Composed of a finite number of states
  - One is designated as the starting state
  - One, or more, are designated as accepting (final) states
- Responds to events
  - Each state responds to some of the events
  - A response “consumes” the event, and may also change the state it is in (moving to “next” state).
- Non-determinism
  - A state responding to an event may have a choice between different states to which it can change.
    - If all states have no such choice then we have a Deterministic Finite Automata (DFA)
    - If any state has such a choice, then we have a Non-deterministic Finite Automata (NDFA)



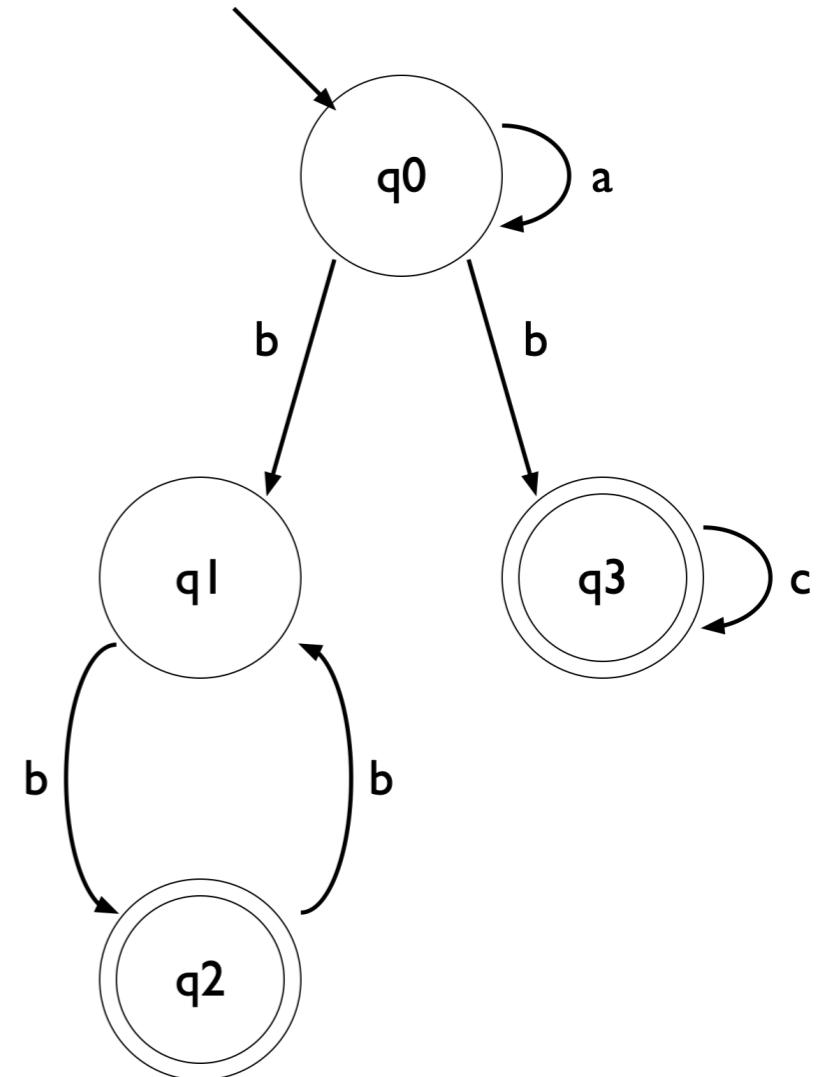
# FSAs “accept” event-sequences

- Given a (N)DFA and a sequence of events, we can ask
  - Can the (N)DFA perform this sequence?
  - Does the starting event “respond” to the first event?
  - If so, does the next state “respond” to the second event?
  - and so on.
- We say that a (N)DFA “accepts” a sequence of events if there is a way to respond to all of the events such that it ends up in one of its “accepting” states.



# NDFA Example

- Regular expressions are a shorthand way to describe finite sequences
  - seq1 seq2
    - seq1 followed by seq2
  - seq1 + seq2
    - either seq1 or seq2
  - seq<sup>\*</sup>
    - zero or more repetitions of seq
  - seq<sup>+</sup>
    - one or more repetitions of seq
- E.g.:  $a^*((bb)^+ + bc^*)$ 
  - zero or more **a**s,  
followed by either:  
an even non-zero number of **b**s;  
or one **b** followed by zero or more **c**s.

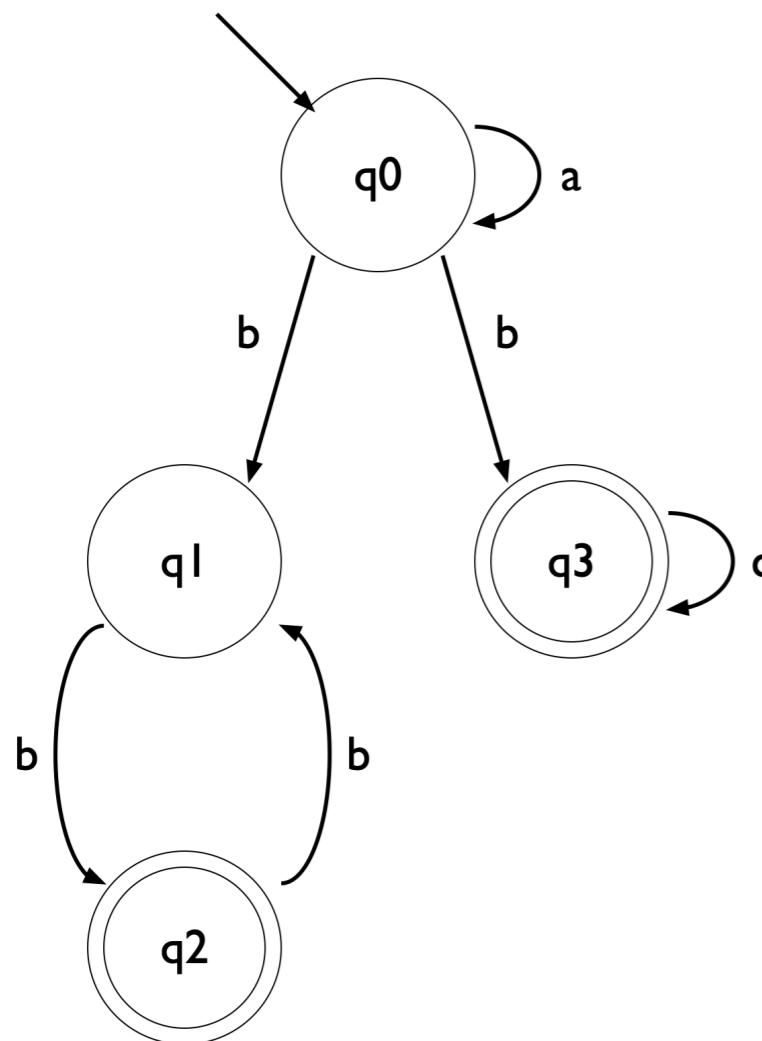


Principles of the Spin Model Checker, M. Ben-Ari, Sec 8.1, pp 125-7



# NDFA in Promela

Any NDFA can be written in Promela



```
1 /* Copyright 2007 by Moti Ben-Ari
2      under the GNU GPL; see readme.txt */
3
4 #define LEN 5
5
6 active proctype FA() {
7     byte h;
8     byte i[LEN];
9     i[0] = 'a'; i[1] = 'a'; i[2] = 'b'; i[3] = 'b'; i[4] = '.';
10    q0: if
11        :: i[h] == 'a' -> h++; goto q0
12        :: i[h] == 'b' -> h++; goto q3
13        :: i[h] == 'b' -> h++; goto q1
14    fi;
15    q1: if
16        :: i[h] == 'b' -> h++; goto q2
17    fi;
18    q2: if
19        :: i[h] == 'b' -> h++; goto q1
20        :: i[h] == '.' -> goto accept
21    fi;
22    q3: if
23        :: i[h] == 'c' -> h++; goto q3
24        :: i[h] == '.' -> goto accept
25    fi;
26    accept:
27        printf("Accepted!\n");
28        // assert(false)
29 }
```



# Applying SPIN to NDFA

- We can simulate the example,
  - `spin SMC-8-1.pml`

but because the first b results in either q1 or q3 being chosen non-deterministically, we are not guaranteed to get the accepting outcome
- We can verify it
  - `spin -run SMC-8-1.pml`

Here we have only 1 process, so we get a “timeout” rather than any deadlock - not helpful.
- We can “falsify” it.
  - We want to see if we can get to **Accepted!** message, so `assert (false)` just afterwards !?!?
  - A verification run will find the only “failure” in the system (`false`), and generate the sequence of steps leading through label “`accept:`”



# The “flip side” of counter-examples

- SPIN is designed to exhaustively check a model to see if any desired property is in fact false.
  - If it finds such a failure, it generates a counter-example - sequence of moves leading to failure. If it doesn't, it reports that all is well.
- We can use SPIN to find solutions to problems too.
  - We design our model so that:
    - (a) there is a point in the code that corresponds to having a solution, or  
(b) we specify a property that is true when a solution is found.
  - We then:
    - (a) add in an **assert(false)** after the solution point in the code, or  
(b) specify the negation of the property
  - SPIN will then, eventually, find the **assert(false)** , or where the negated property is false, and report this as a failure with a “counter-example” that leads to that point - the solution !

